

PA3 – Tabular Q-Learning (Nim)

Due: May 28, 2026 at 11:59 PM

Overview

In this assignment you will implement tabular Q-learning for Nim, a two-player turn-based strategy game. You will train an RL policy through simulated play against simple opponents, then evaluate the learned greedy policy. All graded code tasks are marked with TODO comments in `PA3.ipynb`. Question 4 is graded from the written response prompts in the notebook, using the reporting cell immediately before those prompts.

Learning Goals

- Represent game states and legal actions for Nim.
- Implement tabular Q-learning updates and epsilon-greedy action selection.
- Build an episode-level training loop against random and rule-based opponents.
- Evaluate a learned greedy policy over many matches.
- Interpret a Q-table, a greedy action, and a reported win rate.

Rules and Allowed Libraries

- Use only the packages already imported in the notebook.
- Do not use external RL libraries, such as Gymnasium or Stable-Baselines3.
- All RL logic must be implemented from scratch using Python standard libraries.
- GenAI tools that generate output may be used in course assignments (not tests or exams) to support your learning, as long as you do so honestly through proper documentation, citation, and acknowledgement.

Deliverables

- Submit `PA3.ipynb` with all TODOs in Questions 0–3 completed and the Question 4 written responses filled in.
- Submit the notebook to Gradescope as an `.ipynb` file, not as a PDF.

- Do not rename required functions or change function signatures.
- Optional Question 5 may remain unimplemented.

Point Distribution (100 total)

Section	Points
Question 0: Nim setup and utilities	15
Question 1: Core Q-learning functions	40
Question 2: Training loop	25
Question 3: Policy and evaluation	15
Question 4: Short written responses (manually graded)	5

Optional Extension

Question 5, the SARSA extension, is optional and ungraded. It is included for extra practice comparing an on-policy update with the Q-learning update.

Notes

- The autograded portion of PA3 is worth 95 points total.
- Question 4 is graded manually from the written responses in the notebook.
- Gradescope will show some autograder checks immediately after submission, and additional autograder checks also run during grading but are not shown in advance.
- Reward is always from player 0's perspective: +1 win, -1 loss, 0 otherwise.
- Player 0 always moves first.
- The Question 2 checkpoint trains `Q_small` with `episodes=200`, `alpha=0.5`, `gamma=0.9`, and `epsilon=0.2`. Use that checkpoint output for Question 4.
- For extra debugging after the checkpoint works, you can also try a longer run such as `episodes=10000`, `alpha=0.5`, `gamma=0.9`, `epsilon=0.1`.

Submission Checklist

- All required TODOs in Questions 0–3 are implemented.
- Question 4 written responses are filled in.
- All graded cells through Question 4 run without errors.
- You did not rename required functions or change function signatures.
- Your `.ipynb` file is uploaded to Gradescope and the autograder shows no submission errors.
- Optional Question 5 may remain unimplemented.

Tips

- Complete Question 0 first; most later functions depend on these utilities.
- In Question 2, a nonterminal Q-learning transition includes the agent move and the opponent response. Update the agent's (`state`, `action`) using the state after the opponent moves.
- Run each nearby checkpoint cell after finishing a TODO.
- You may submit multiple times; use early Gradescope feedback to catch interface errors before working on later questions.